




1. State some autoencoders and PCA.
2. What is the back-propagation algorithm for an autoencoder layer?
3. Assume that you initialize all weights in a neural net to the same value and you do the same for the bias terms. Is this a good idea? Justify your answer.
4. An autoencoder is a neural network designed to learn feature representations in an unsupervised manner. Unlike a standard multi-layer network, an autoencoder has the same number of nodes in its output layer as its input layer. An autoencoder is trained to reconstruct its own input x , i.e. to minimize the reconstruction error. An autoencoder is shown below.



Suppose the input is a set of P -dimensional unlabelled data $\{x^{(i)}\}_{i=1}^N$. Consider an autoencoder with H hidden units in the second layer L_2 . We will use the following notation for this autoencoder:

- W^e denotes the $P \times H$ weight matrix between L_1 and L_2
- W^d denotes the $H \times P$ weight matrix between L_2 and L_3
- σ denotes the activation function for L_2 and L_3
- $s_j^{(i)} = \sum_{k=1}^P W_{kj}^e x_k^{(i)}$
- $h_j^{(i)} = \sigma(\sum_{k=1}^P W_{kj}^e x_k^{(i)})$
- $t_j^{(i)} = \sum_{k=1}^H W_{kj}^d h_k^{(i)}$
- $\hat{x}_j^{(i)} = \sigma(\sum_{k=1}^H W_{kj}^d h_k^{(i)})$

- $J(W^e, W^d) = \sum_{j=1}^N \frac{1}{2} \|x^{(i)} - \hat{x}^{(i)}\|_2^2 = \sum_{j=1}^N \frac{1}{2} (x_j^{(i)} - \hat{x}_j^{(i)})^2$ is the reconstruction error for example $x^{(i)}$
- $J(W^e, W^d) = \sum_{j=1}^N J(W^e, W^d)^{(i)}$ is the total reconstruction error
- (We add an input layer and hidden layer so that no bias term has to be added)



Fill in the missing terms in the equations for W^e and W^d . Use the notation defined above. No new notation needed.

$$\frac{\partial J}{\partial W_{kl}^d} = \sum_{j=1}^N \left(\boxed{} \cdot \frac{\partial \hat{x}_j^{(i)}}{\partial W_{kl}^d} \right)$$

$$\frac{\partial \hat{x}_j^{(i)}}{\partial W_{kl}^d} = \sigma \left(\sum_{k=1}^H W_{kj}^e x_k^{(i)} \right) \cdot \boxed{}$$


$$\frac{\partial J}{\partial W_{kl}^e} = \frac{\partial J^{(i)}}{\partial s_j^{(i)}} \cdot \boxed{}$$

$$\frac{\partial J^{(i)}}{\partial s_j^{(i)}} = \sum_{k=1}^H \left(\frac{\partial J^{(i)}}{\partial x_k^{(i)}} \cdot \boxed{} \cdot \sigma'(s_j^{(i)}) \right)$$

QQ: 749389476

5. 3σ rule is a common technique used for anomaly detection. Describe what is the intuition of this rule for anomaly detection? How our result will be effected if we use other values of σ (e.g., 2σ , or 4σ)?

<https://tutorcs.com>



6. In the VAE, how sampling of the latent code is different during training and generation (generating a new sample)?